
Image Compression

1. What and why image compression
2. Basic concepts
3. Encoding/decoding, entropy

What is Data and Image Compression?

- Data compression is the art and science of representing information in a compact form.
- Data is a sequence of symbols taken from a discrete alphabet.

Why do we need Image Compression?

Still Image

- One page of A4 format at 600 dpi is > 100 MB.
- One color image in digital camera generates 10-30 MB.
- Scanned 3"×7" photograph at 300 dpi is 30 MB.

Digital Cinema

- $4K \times 2K \times 3 \times 12$ bits/pel = 48 MB/frame or 1 GB/sec
or 70 GB/min.

Why do we need Image Compression?

- 1) Storage
- 2) Transmission
- 3) Data access

1990-2000

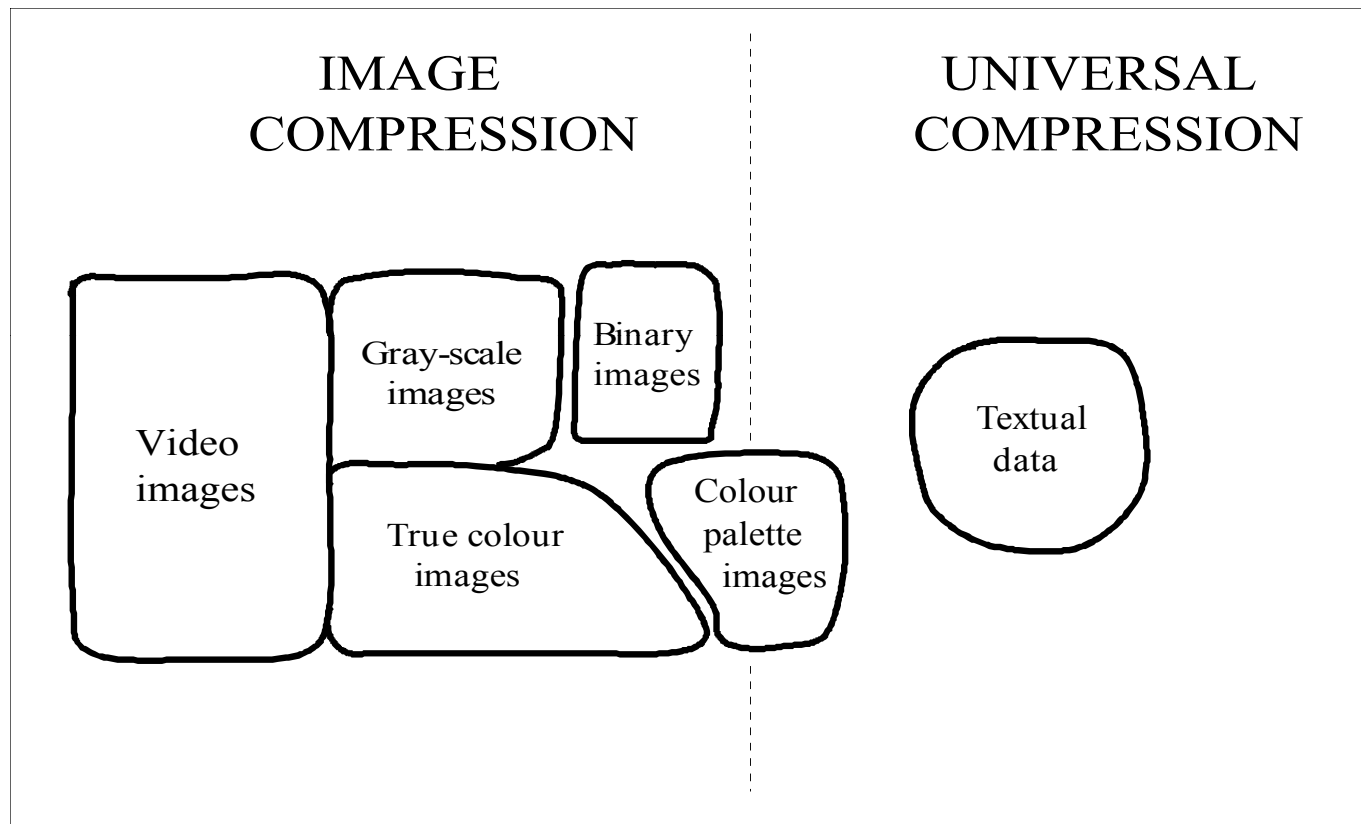
Disc capacities : 100MB -> 20 GB (200 times!)
but seek time : 15 milliseconds → 10 milliseconds
and transfer rate : 1MB/sec -> 2 MB/sec.

Compression improves overall response time in some applications.

Source of images

- Image scanner
- Digital camera
- Video camera,
- Ultra-sound (US), Computer Tomography (CT),
Magnetic resonance image (MRI), digital X-ray (XR),
Infrared.
- etc.*

Image types



Why we can compress image?

- Statistical redundancy:
 - 1) Spatial correlation
 - a) Local - Pixels at neighboring locations have similar intensities.
 - b) Global - Reoccurring patterns.
 - 2) Spectral correlation – between color planes.
 - 3) Temporal correlation – between consecutive frames.
- Tolerance to fidelity:
 - 1) Perceptual redundancy.
 - 2) Limitation of rendering hardware.

Lossy vs. Lossless compression

Lossless compression: reversible, information preserving
text compression algorithms,
binary images, palette images

Lossy compression: irreversible
grayscale, color, video

Near-lossless compression:
medical imaging, remote sensing.

Rate measures

Bitrate: $\frac{\text{size of the compressed file}}{\text{pixels in the image}} = \frac{C}{N} \quad \text{bits/pel}$

Compression ratio: $\frac{\text{size of the original file}}{\text{size of the compressed file}} = \frac{N \cdot k}{C}$

Distortion measures

Mean average error (MAE):
$$\text{MAE} = \frac{1}{N} \sum_{i=1}^N |y_i - x_i|$$

Mean square error (MSE):
$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - x_i)^2$$

Signal-to-noise ratio (SNR):
$$\text{SNR} = 10 \cdot \log_{10} [\sigma^2 / \text{MSE}]$$

(decibels)

Pulse-signal-to-noise ratio (PSNR):
$$\text{PSNR} = 10 \cdot \log_{10} [A^2 / \text{MSE}]$$

(decibels)

A is amplitude of the signal: $A = 2^8 - 1 = 255$ for 8-bits signal.

Other issues

- Coder and decoder computation complexity
- Memory requirements
- Fixed rate or variable rate
- Error resilience
- Symmetric or asymmetric
- Decompress at multiple resolutions
- Decompress at various bit rates
- Standard or proprietary

Entropy

Set of symbols (alphabet) $S=\{s_1, s_2, \dots, s_N\}$,

N is number of symbols in the alphabet.

Probability distribution of the symbols: $P=\{p_1, p_2, \dots, p_N\}$

According to Shannon, the entropy H of an information source S is defined as follows:

$$H = -\sum_{i=1}^N p_i \cdot \log_2(p_i)$$

Entropy

The amount of information in symbol s_i ,
in other words, the number of bits to code or code length
for the symbol s_i :

$$H(s_i) = -\log_2(p_i)$$

The average number of bits for the source S:


$$H = -\sum_{i=1}^N p_i \cdot \log_2(p_i)$$

Entropy for binary source: $N=2$

$$S=\{0,1\}$$

$$p_0=p$$

$$p_1=1-p$$




$$H = -(p \cdot \log_2 p + (1-p) \cdot \log_2 (1-p))$$

$$H=1 \text{ bit for } p_0=p_1=0.5$$

Entropy for uniform distribution: $p_i=1/N$

Uniform distribution of probabilities: $p_i=1/N$:

$$H = -\sum_{i=1}^N (1/N) \cdot \log_2(1/N) = \log_2(N)$$



Examples:

$$N=2: \quad p_i=0.5; \quad H=\log_2(2) = 1 \text{ bit}$$

$$N=256: \quad p_i=1/256; \quad H=\log_2(256)= 8 \text{ bits}$$

How to get the probability distribution?

1) Static modeling:

- a) The same code table is applied to all input data.
- b) One-pass method (encoding)
- c) No side information

2) Semi-adaptive modeling:

- a) Two-pass method: (1) analysis and (2) encoding.
- b) Side information needed (model, code table)

3) Adaptive (dynamic) modeling:

- a) One-pass method: analysis and encoding
- b) Updating the model during encoding/decoding
- c) No side information

Static vs. Dynamic: Example

$S = \{a,b,c\}$; Data: a,a,b,a,a,c,a,a,b,a.

1) Static model: $p_i = 1/10$

$$H = -\log_2(1/10) = \mathbf{1.58} \text{ bits}$$

2) Semi-adaptive method: $p_1 = 7/10$; $p_2 = 2/10$; $p_3 = 1/10$;

$$H = -(0.7 \cdot \log_2 0.7 + 0.2 \cdot \log_2 0.2 + 0.1 \cdot \log_2 0.1) = \mathbf{1.16} \text{ bits}$$

3) Adaptive method: Example

$S = \{a,b,c\}$; Data: a,a,b,a,a,c,a,a,b,a.

Symbol	1	2	3	4	5	6	7	8	9	10
a	1	2	3	3	4	5	5	6	7	7
b	1	1	1	2	2	2	2	2	2	3
c	1	1	1	1	1	1	2	2	2	2
p_i	0.33	0.5	0.2	0.5	0.57	0.13	0.56	0.60	0.18	0.58
H	1.58	1.0	2.32	1.0	0.81	3.0	0.85	0.74	2.46	0.78

$$H = (1/10)(1.58 + 1.0 + 2.32 + 1.0 + 0.81 + 3.0 + 0.85 + 0.74 + 2.46 + 0.78) \\ = \mathbf{1.45} \text{ bits/char}$$

$$1.16 < 1.45 < 1.58$$

S.-Ad. Ad. Static

Coding methods

- Shannon-Fano Coding
 - Huffman Coding
 - Predictive coding
 - Block coding
-
- Arithmetic code
 - Golomb-Rice codes

Shannon-Fano Code: A top-down approach


1) Sort symbols according their probabilities:

$$p_1 \leq p_2 \leq \dots \leq p_N$$

2) Recursively divide into parts, each with approx. the same number of counts (probability)

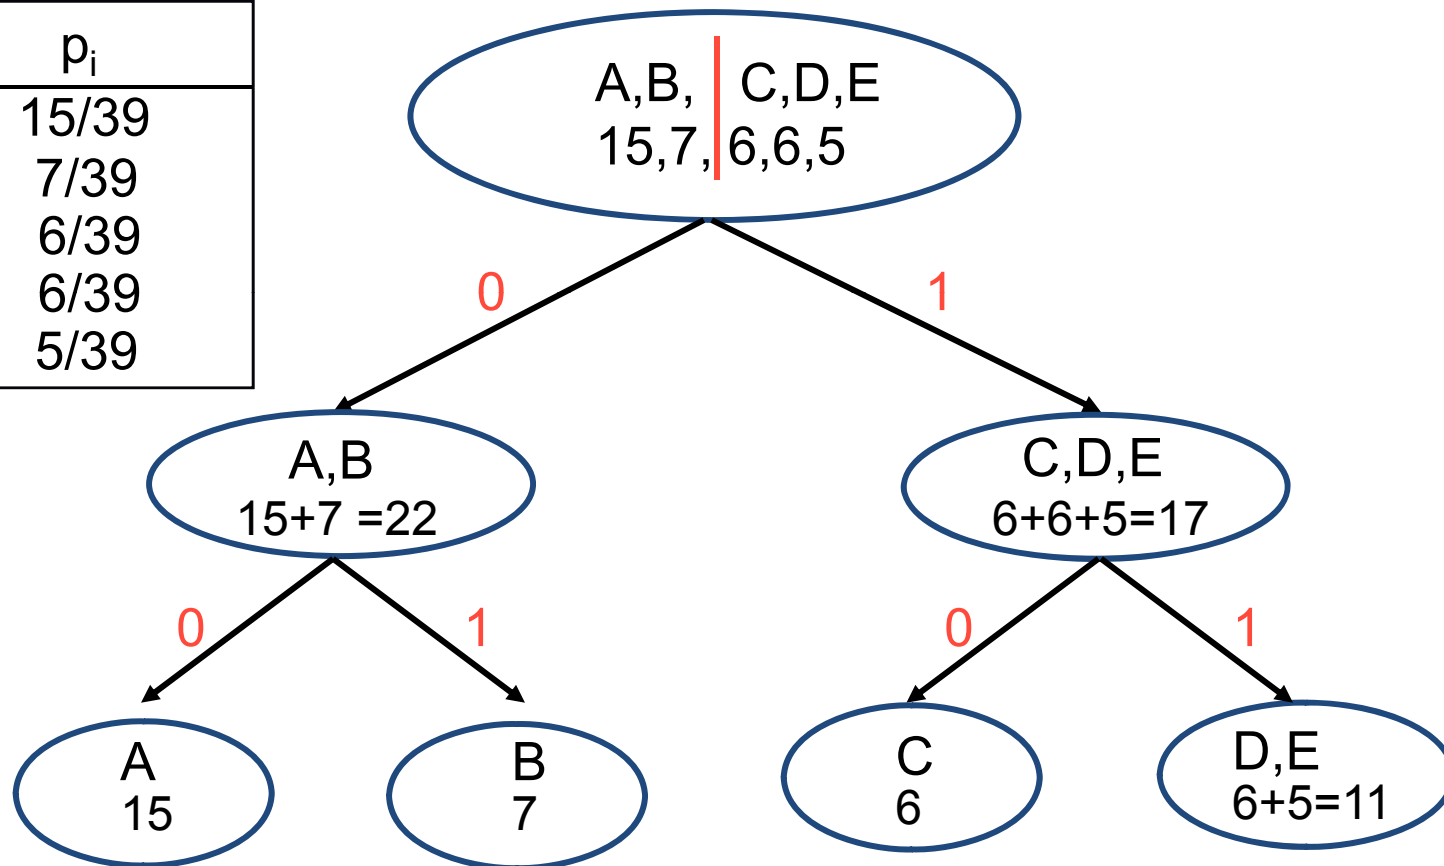
Shannon-Fano Code: Example (1 step)

s_i	p_i
A -	15/39
B -	7/39
C -	6/39
D -	6/39
E -	5/39



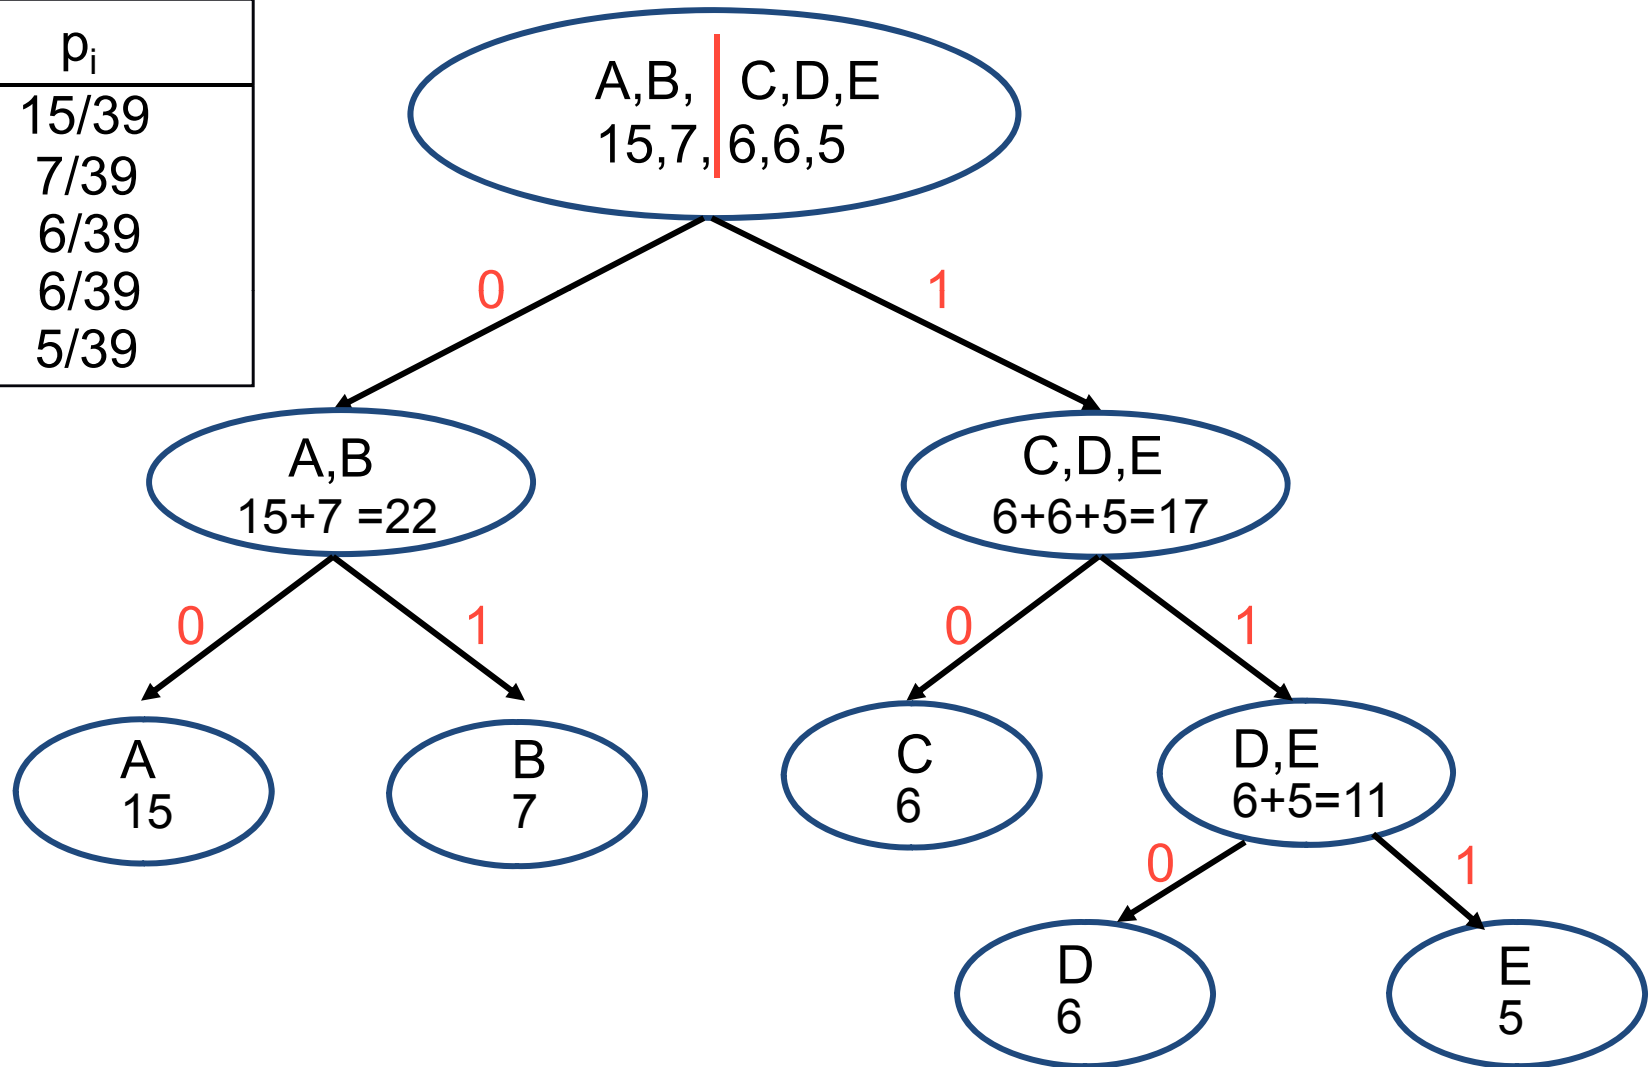
Shannon-Fano Code: Example (2 step)

s_i	p_i
A -	15/39
B -	7/39
C -	6/39
D -	6/39
E -	5/39




Shannon-Fano Code: Example (3 step)

s_i	p_i
A -	15/39
B -	7/39
C -	6/39
D -	6/39
E -	5/39



Shannon-Fano Code: Example (Result)

Symbol	p_i	$-\log_2(p_i)$	Code	Subtotal
A	15/39	1.38	00	2×15
B	7/39	2.48	01	2×7
C	6/39	2.70	10	2×6
D	6/39	2.70	110	3×6
E	5/39	2.96	111	3×5
Total:				89 bits



Binary tree

$$H = 89/39 = 2.28 \text{ bits}$$


Shannon-Fano Code: Encoding

A - 00
B - 01
C - 10
D - 110
E - 111

Message: B A B A C A C A D E

Codes: 01 00 01 00 10 00 10 00 110 111

Bitstream: 0100010010001000110111



Binary tree


Shannon-Fano Code: Decoding

A - 00
B - 01
C - 10
D - 110
E - 111

Bitstream: 0100010010001000110111 (23 bits)

Codes: 01 00 01 00 10 00 10 00 110 111

Message: B A B A C A C A D E



Binary tree

Huffman Code: A bottom-up approach

INIT:


Put all nodes in an OPEN list, keep it sorted all times according their probabilities;.

REPEAT

- a) From OPEN pick *two* nodes having the *lowest* probabilities, create a parent node of them.
- b) Assign the sum of the children's probabilities to the parent node and inset it into OPEN
- c) Assign code 0 and 1 to the two branches of the tree, and delete the children from OPEN.

Huffman Code: Example

Symbol	p_i	$-\log_2(p_i)$	Code	Subtotal
A	15/39	1.38	0	2×15
B	7/39	2.48	100	3×7
C	6/39	2.70	101	3×6
D	6/39	2.70	110	3×6
E	5/39	2.96	111	3×5
Total:				87 bits



Binary tree

$$H = 87/39 = 2.23 \text{ bits}$$


Huffman Code: Decoding

A - 0
B - 100
C - 101
D - 110
E - 111

Bitstream: 1000100010101010110111 (22 bits)

Codes: 100 0 100 0 101 0 101 0 110 111

Message: B A B A C A C A D E




Binary tree

Properties of Huffman code

- Optimum code for a given data set requires two passes.
- Code construction complexity $O(N \log N)$.
- Fast lookup table based implementation.
- Requires at least one bit per symbol.
- Average codeword length is within one bit of zero-order entropy (Tighter bounds are known): $H \leq R < H+1$ bit
- Susceptible to bit errors.

Unique prefix property

No code is a prefix to any other code, all symbols are the *leaf* nodes



Shannon-Fano and Huffman codes are prefix codes

Legend: Shannon (1948) and Fano (1949);

Huffman (1952) was student of Fano at MIT.

Fano: "Construct minimum-redundancy code → final exam is passed!"

Predictive coding




- 1) Calculate prediction value: $y_i = f(\text{neighbourhood of } x_i)$.
- 2) Calculating the prediction error: $e_i = y_i - x_i$.
- 3) Encode the prediction error e_i .

Predictive model for grayscale images



$$y = x_i - x_{i-1}$$

.....→




Histogram of the original image and Residual image

Entropy: $H_o = 7.8$ bits/pel (?)

$H_r = 5.1$ bits/pel (?)

Coding without prediction




● $f_0=8$; $p_0=p=8/64 =0.125$;

○ $f_1=56$; $p_1=(1-p)=56/64=0.875$

Entropy:

$$H = -((8/64)*\log_2(8/64)+(56/64)*\log_2(56/64))=\mathbf{0.544} \text{ bits/pel}$$

Prediction for binary images by pixel above




Entropy:

$$H = -((16/64) \cdot \log_2(16/64) + (48/64) \cdot \log_2(48/64)) = \mathbf{0.811} \text{ bits/pel}$$

Wrong predictor!

Prediction for binary images pixel to the left



Entropy:

$$H = -((1/64) \cdot \log_2(1/64) + (63/64) \cdot \log_2(63/64)) = \mathbf{0.116} \text{ bits/pel}$$

Good predictor!

Comparison of predictors:

- Without prediction: $H = 0.544$ bits/pel
- Prediction by pixel above: $H = 0.811$ bits /pel (bad!)
- Prediction by pixel to the left: $H = \mathbf{0.116}$ bits/pel (good!)